Quarterly Technology Report. Retrieval-Augmented Generation, RAG, architecture overview. The hybrid retrieval pipeline combines BM25 keyword search with dense vector embeddings from BAAI/bge-m3.

Model selection. The team standardised on NVIDIA RTX-4090 GPUs with 128GB memory for local inference benchmarks.